

1. Generate random addresses with the following arguments: -s 0 -n 10, -s 1 -n 10, and -s 2 -n 10. Change the policy from FIFO, to LRU, to OPT. Compute whether each access in said address traces are hits or misses.
→ See Ordner
2. For a cache of size 5, generate worst-case address reference streams for each of the following policies: FIFO, LRU, and MRU (worst-case reference streams cause the most misses possible. For the worst case reference streams, how much bigger of a cache is needed to improve performance dramatically and approach OPT?

Worst case FIFO:

```
./paging-policy.py -a  
0,1,2,3,4,5,6,7,8,9,1,2,3,4,5,6,7,8,9,0,1,2,3,4,5,6,7,8,9,1,2,3,4,5,6,7,8,9 -c  
-p FIFO -C 5
```

- to approach OPT the cache size must be increased by 5, so it is 10.
- Number of unique pages = 10

Worst case LRU:

```
./paging-policy.py -a  
0,1,2,3,4,5,6,7,8,9,1,2,3,4,5,6,7,8,9,0,1,2,3,4,5,6,7,8,9,1,2,3,4,5,6,7,8,9 -c  
-p LRU -C 5
```

- to approach OPT the cache size must be increased by 5, so it is 10.
- Number of unique pages = 10

Worst case MRU:

```
./paging-policy.py -a 0,1,2,3,4,5,4,5,4,5,4,5 -c -p MRU -C 5
```

- Worst case when you have (cache size + 1) unique pages and then you oscillate between two

- Increasing the cache size by one solves the problem
- New cache size is 6, which approaches OPT

3. Generate a random trace (use python or perl). How would you expect the different policies to perform on such a trace?

- run generate-random.py

```
./paging-policy.py -a 3, 1, 6, 0, 5, 3, 0, 5, 0, 4 -c -p FIFO -C 5
./paging-policy.py -a 3, 1, 6, 0, 5, 3, 0, 5, 0, 4 -c -p LRU -C 5
./paging-policy.py -a 3, 1, 6, 0, 5, 3, 0, 5, 0, 4 -c -p RAND -C 5
./paging-policy.py -a 3, 1, 6, 0, 5, 3, 0, 5, 0, 4 -c -p MRU -C 5
./paging-policy.py -a 3, 1, 6, 0, 5, 3, 0, 5, 0, 4 -c -p OPT -C 5
```

- I expect all of them to perform roughly the same. With OPT being the best
- They all get a hit rate of 40%

```
./paging-policy.py -s 0 -n 10 -c -p FIFO
FINALSTATS hits 1   misses 9   hitrate 10.00

./paging-policy.py -s 0 -n 10 -c -p LRU
FINALSTATS hits 2   misses 8   hitrate 20.00

./paging-policy.py -s 0 -n 10 -c -p RAND
FINALSTATS hits 0   misses 10   hitrate 0.00

./paging-policy.py -s 0 -n 10 -c -p OPT
FINALSTATS hits 4   misses 6   hitrate 40.00

./paging-policy.py -s 0 -n 10 -c -p UNOPT
FINALSTATS hits 0   misses 10   hitrate 0.00

./paging-policy.py -s 0 -n 10 -c -p MRU
FINALSTATS hits 2   misses 8   hitrate 20.00
```

4. Now generate a trace with some locality. How can you generate such a trace? How does LRU perform on it? How much better than RAND is LRU? How does CLOCK do? How about CLOCK with different numbers of clock bits?

→ run generate-locality.py

```
./paging-policy.py -a 9,9,1,9,5,5,9,9,8,9 -c -p FIFO  
FINALSTATS hits 5   misses 5   hitrate 50.00
```

```
./paging-policy.py -a 9,9,1,9,5,5,9,9,8,9 -c -p LRU  
FINALSTATS hits 6   misses 4   hitrate 60.00
```

```
./paging-policy.py -a 9,9,1,9,5,5,9,9,8,9 -c -p RAND  
FINALSTATS hits 6   misses 4   hitrate 60.00
```

```
./paging-policy.py -a 9,9,1,9,5,5,9,9,8,9 -c -p CLOCK -b 0  
FINALSTATS hits 6   misses 4   hitrate 60.00
```

```
./paging-policy.py -a 9,9,1,9,5,5,9,9,8,9 -c -p CLOCK -b 1  
FINALSTATS hits 6   misses 4   hitrate 60.00
```

```
./paging-policy.py -a 9,9,1,9,5,5,9,9,8,9 -c -p CLOCK -b 2  
FINALSTATS hits 6   misses 4   hitrate 60.00
```

5. Use a program like valgrind to instrument a real application and generate a virtual page reference stream. For example, running `valgrind --tool=lackey --trace-mem=yes ls` will output a nearly-complete reference trace of every instruction and data reference made by the program `ls`. To make this useful for the simulator above, you'll have to first transform each virtual memory reference into a virtual page-number reference (done by masking off the offset and shifting the resulting bits downward). How big of a cache is needed for your application trace in order to satisfy a large fraction of requests? Plot a graph of its working set as the size of the cache increases.

```
valgrind --tool=lackey --trace-mem=yes ls &> ref-trace.txt
```

→ Wenn man `ref-trace.txt` analysiert, sieht man dass die ersten 5 hex Zahlen zur VPN gehören, und die letzten 3 zum offset

